

# Wine quality score prediction through reviews sentiment analysis and regression models Notebook

March 1, 2021

```
[1]: import pandas as pd
import numpy as np
import nltk
import matplotlib.pyplot as plt
from nltk import word_tokenize
from nltk import pos_tag
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords

from nltk.corpus import wordnet as wn
from nltk.stem.wordnet import WordNetLemmatizer
from nltk import word_tokenize, pos_tag
from collections import defaultdict

from sklearn.feature_extraction import DictVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
from scipy.sparse import hstack
from sklearn.utils import shuffle
from sklearn import metrics
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.neural_network import MLPRegressor
from sklearn.svm import SVR
from sklearn.model_selection import cross_val_score
from wordcloud import WordCloud, STOPWORDS
import time
import seaborn as sns
```

```
[2]: nltk.download('wordnet')
nltk.download('punkt')
nltk.download('stopwords')
```

```
[nltk_data] Downloading package wordnet to
[nltk_data]   /home/bigdata-01QYD/s259444/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]   /home/bigdata-01QYD/s259444/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data]   /home/bigdata-01QYD/s259444/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

[2]: True

```
[3]: column_names=["country", "description", "designation", "province", "region_1", "region_2", "variety", "winery", "quality"]
wine_df = pd.read_csv('dev.tsv', sep='\t', header=1, names = column_names)
wine_df.head()
```

```
[3]:
```

	country	description	designation	
0	US	Simple and dry, this Cabernet has modest black...	NaN	
1	US	This lovely wine captures the floral, perfumed...	NaN	
2	Portugal	The aromas are the thing here, as so often wit...	NaN	
3	Italy	This is an interesting, outright strange wine ...	Nati	
4	New Zealand	Classic gooseberry and pink grapefruit notes f...	NaN	

	province	region_1	region_2	variety	
0	California	Paso Robles	Central Coast	Cabernet Sauvignon	
1	Oregon	Willamette Valley	Willamette Valley	Gewürztraminer	
2	Alentejano	NaN	NaN	Touriga Nacional	
3	Southern Italy	Pompeiano	NaN	Coda di Volpe	
4	Marlborough	NaN	NaN	Sauvignon Blanc	

	winery	quality
0	Castle Rock	31.0
1	Château Bianca	35.0
2	Herdade do Esporão	41.0
3	Sorrentino	37.0
4	Cloudy Bay	48.0

```
[4]: wine_df.shape
```

[4]: (120743, 9)

The dataset has 120743 rows and 9 columns. For the columns we have the country from where the vine has been produced, the province and the two regions. So the province includes the regions, and the country includes the province.

Then there is the designation that is the name give to the wine by the produce, the variety describes the type of grapes that are used, then the winery from which is produced.

Then there's the description provided by the reviewer, and also the quality scores.

All the attributes are categorical and nominal, exception made for the quality one, that is numerical and ordinal too.

```
[5]: wine_df.describe()
```

```
[5]:
```

	quality
count	120743.000000
mean	46.277863
std	11.924830
min	0.000000
25%	38.000000
50%	46.000000
75%	55.000000
max	100.000000

from the table above, we can see the mean and the standard deviation, and also understand partially its distribution thanks to the percentiles.

```
[6]: wine_df.isnull().values.any()
```

```
[6]: True
```


```
[7]: wine_df.isnull().sum()
```

```
[7]:
```

country	5
description	0
designation	36518
province	5
region_1	20008
region_2	72007
variety	0
winery	0
quality	0

dtype: int64

```
[8]: ax = sns.heatmap(wine_df.isnull(), yticklabels=False, cbar=False, cmap='viridis')  
ax.figure.savefig('null_heatmap.png')
```



The table and plot show that region\_2 has more than the half of null values, region\_1 20000 records and designation 30000. There also some records (5) that have also the country value null. So as first solution, designation, region\_1 and region\_2 columns are dropped, and also the 5 rows with country value null.

```
[4]: wine_df = wine_df.dropna(how="any", subset=["winery", "country", "variety", "province"])
      wine_df = wine_df.drop(["designation", "region_1", "region_2"], axis=1)
```

```
[10]: wine_df.shape
```

```
[10]: (120738, 6)
```

Let's now do a brief cardinality analysis of the attributes of the dataset.

```
[11]: for column in wine_df.columns:
      uniq = np.unique(wine_df[column])
      print('{}: {} distinct values\n'.format(column, len(uniq)))
```

```
country: 48 distinct values
```

```
description: 85001 distinct values
```

province: 444 distinct values


variety: 603 distinct values

winery: 14104 distinct values

quality: 86 distinct values

It's kind of funny to notice that only half of the descriptions are unique.

```
[12]: votes = wine_df.groupby(['quality']).size()
indexes = votes.index
fig, ax = plt.subplots(figsize=(10,7))
ax.bar(indexes, votes)
plt.show()
fig.savefig('score_distribution.png')
```




This is the distribution of the quality score. It's very close to a normal one, with 46,27 as mean.

Now just out of curiosity, let's investigate on how many zeros there are in the quality scores. Because if someone gives 0 in a range of 100 to the scores, the wine must not be wine but winegar. So we can consider them like kinda of outlier


```
[13]: print("There are {} evaluations were the score assigned is 0" .format(votes.
      ↪iloc[0]) )
```

There are 15 evaluations were the score assigned is 0


```
[19]: val_for_country = wine_df.groupby('country').size()
plt.figure(figsize=(10,13))
ax = sns.barplot(x=val_for_country, y=val_for_country.index)
plt.title('# of records for country')
plt.xlabel('n_val')
plt.ylabel('Country')
ax.figure.savefig('records_for_country.png')
```



```
[23]: country_avg = wine_df.groupby('country').mean()['quality'].
      ↪sort_values(ascending=False).to_frame()
plt.figure(figsize=(15,11))
ax = sns.barplot(x=country_avg['quality'], y=country_avg.index)
plt.title('Country average scores')
plt.xlabel('scores')
plt.ylabel('Country')
ax.figure.savefig('country_average_scores.png')
```



```
[16]: wine_df['descrln'] = wine_df['description'].apply(lambda descr:len(descr.
      ↪split()))
len_series = wine_df.groupby('descrln').size()
fig, ax = plt.subplots (figsize = (10,7))
ax.bar(len_series.index, len_series)
ax.set_xlabel('description lenght')
ax.set_ylabel('occurencies')
plt.show()
plt.savefig('descrln_distribution.png')
```




<Figure size 432x288 with 0 Axes>

The plot shows the distribution of the description lengths, to be interpreted as number of words for description. This one too is a normal distribution, with averagea around 40.

```
[17]: sns.set_theme(style='ticks')
      ax = sns.jointplot(x=wine_df['quality'], y=wine_df['descrlen'], kind='hex',
      ↪color="#4CB391")
      ax.fig.savefig('descrlen-quality_jointplot.png')
```





Now, let's divide the scores in 5 equal ascending categories. For computing the division we use percentiles. So we have five percentiles and for the grades: Average one, then descending low grade rated, and bad rated, while ascending from the average grade, we have discrete and great rated. So let's plot the word clouds just to make an idea of what the descriptions tell for the different grades.

```
[18]: ix = [p for p in range(20,120,20)]
perc = [np.percentile(wine_df['quality'], p) for p in range(20,120,20)]
categories=['bad', 'low-grade', 'average', 'discrete', 'great']
sns.set_context("talk")
for i,p in enumerate(perc):
    print("")
    plt.figure(figsize= (16,8))
```



[illegible]

11

[illegible]

Word cloud of Description of great rated wine

```
[5]: wnl = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))
```

```

tag_map = defaultdict(lambda : wn.NOUN)
tag_map['J'] = wn.ADJ
tag_map['V'] = wn.VERB
tag_map['R'] = wn.ADV

def tolemmas(text):
    tokens = word_tokenize(text)
    lemma_function = WordNetLemmatizer()
    lemmas = ""
    for token, tag in pos_tag(tokens):
        lemma = lemma_function.lemmatize(token, tag_map[tag[0]])
        #if not lemma in stop_words:
        lemmas = lemmas + lemma + " "
    return lemmas

```

```

[6]: wine_df['description'] = wine_df['description'].replace('[^a-zA-Z0-9]', ' ', _
    ↪ regex = True)
wine_df['description'] = wine_df['description'].str.lower()
wine_df['description'] = wine_df['description'].map(lambda sentence:
    ↪ tolemmas(sentence))

```

In the two cells above starts the preprocessing of the description attribute. The categorical attributes cannot be given to the machine learning model as they are. Because those models only works with numbers. So first thing first, punctuation marks are removed, then all the text is converted to lower letters. The next important step is to lemmatize words. From wikipedia, lemmatisation is the process of grouping together the inflected forms of a word so they can be analysed as a single item, identified by the word's lemma, or dictionary form. Unlike stemming, lemmatisation depends on correctly identifying the intended part of speech and meaning of a word in a sentence, as well as within the larger context surrounding that sentence, such as neighboring sentences or even an entire document. This is why, in the function defined above, nouns, verbs, adverbs and pronoun are mapped. So some examples can be was->be mice->mouse ->meeting->meet. This is different from stemming that is the process of reducing inflected words to their word stem, base or root form—generally a written word form.

After doing this, each description, that is a vector of words, is encoded with Tfidf. Deeply analyzed during the course, Tfidf is a numerical statistic that is intended to reflect how important a word is to a document in a collection or corpus. The tf-idf value increases proportionally to the number of times a word appears in the document and is offset by the number of documents in the corpus that contain the word, which helps to adjust for the fact that some words appear more frequently in general. One important thing to be aware is that the model with this encoding does not learn the “meaning” of the words, for that there are other types of encoding and models.

After doing different simulations, I decided to not remove stop words, because they can bring additional value to the model, in fact the r2 score is higher with stopword than without. This I think is more valid also for a different model, that I will not explore in this notebook, because was only mentioned in theory lessons of the course, but not deeply analyzed.

```
[7]: vect_tfidf = TfidfVectorizer(min_df=5)
descrtfidf = vect_tfidf.fit_transform(wine_df['description']).toarray()

vect_1hot = DictVectorizer(sparse=False, dtype=int)
wine_df_tohot = wine_df.filter(['country', 'province', 'variety', 'winery'],
    ↪axis=1)
data = wine_df_tohot.to_dict('records')
cat_1hot = vect_1hot.fit_transform(data)

X_dev = np.concatenate([descrtfidf, cat_1hot], axis=1)
y_dev = wine_df['quality']
X_dev, y_dev = shuffle(X_dev, y_dev, random_state=0)
```

For the other categorical attributes, one-hot encoding is used.

For the development dataset in both tfidf and one-hot encoding fit\_transform is used, so this means that the distribution of the data is fitted and then the same data is transformed to the encoding

In the cell above, the evaluation dataset is loaded, and the same preprocessing steps are made. Here no rows are dropped, cause these missing lines would obviously cause issues with the submission. The difference with the development set is that here for both tfidf and one-hot encoding we use first the fit method and then the transform one. This because the dev set is way more larger than the evaluation one, so “reality” can be better represented.

```
[8]: column_names_ev=["country", "description", "designation", "province",
    ↪"region_1", "region_2", "variety", "winery"]
wine_df_ev = pd.read_csv('eval.tsv', sep='\t', header=0, names =
    ↪column_names_ev)
wine_df_ev = wine_df_ev.drop(["designation", "region_1", "region_2"], axis=1)

wine_df_ev['description'] = wine_df_ev['description'].replace('[^a-zA-Z0-9]', '
    ↪', regex = True)
wine_df_ev['description'] = wine_df_ev['description'].str.lower()
wine_df_ev['description'] = wine_df_ev['description'].map(lambda sentence:
    ↪tolemmas(sentence))

vect_tfidf = TfidfVectorizer(min_df=5)
vect_tfidf.fit(wine_df['description'])
descr_ev_tfidf = vect_tfidf.transform(wine_df_ev['description']).toarray()

wine_df_ev_tohot = wine_df_ev.filter(['country', 'province', 'variety',
    ↪'winery'], axis=1)
data_ev = wine_df_ev_tohot.to_dict('records')
vect_1hot = DictVectorizer(sparse=False, dtype=int)
vect_1hot.fit(data)
cat_1hot_ev = vect_1hot.transform(data_ev)
```



```
X_ev = np.concatenate([descr_ev_tfidf, cat_1hot_ev], axis=1)
```

```
[23]: ridge_params = {'alpha': [0.1, 0.2, 0.7]}
```

```
[24]: clf = Ridge(alpha=0.1)
scores = cross_val_score(clf, X_dev, y_dev, cv=4)
print(scores)
```

```
[0.7107029  0.71915213 0.71570793 0.71326291]
```

```
[25]: clf = Ridge(alpha=0.2)
scores = cross_val_score(clf, X_dev, y_dev, cv=4)
print(scores)
```

```
[0.71754641 0.72590613 0.72302613 0.72092031]
```

Since the r2 scores of the cross validation has no large variance, and since the RAM resources of the account at Jupyter Polito are limited, the other hyperparameter setting of alpha will be explored without cross validation. Also because for now, no PCA for dimensionality reduction is applied.

```
[10]: seed=1234
X_train, X_val, y_train, y_val = train_test_split(X_dev, y_dev, test_size = 0.
↪2, random_state=seed)
```

```
[26]: clf = Ridge(alpha=0.5)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_val)
y_true = y_val
r2 = r2_score(y_true, y_pred)
print("r2 score for alpha=0.5: {}".format(r2))
```

```
r2 score for alpha=0.5: 0.7339390879492909
```

```
[27]: clf = Ridge(alpha=0.7)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_val)
y_true = y_val
r2 = r2_score(y_true, y_pred)
print("r2 score for alpha=0.7: {}".format(r2))
```

```
r2 score for alpha=0.7: 0.7344847198556729
```

```
[28]: clf = Ridge(alpha=0.5)
clf.fit(X_dev, y_dev)
y_pred = clf.predict(X_ev)

pred_df = pd.DataFrame()
pred_df['Id'] = wine_df_ev.index
```

```
pred_df['Predicted'] = y_pred

pred_df.to_csv('Final_sub_ridge_0.5.csv', header=True, index=None)
```

After trying to predict the score with lasso model, it's time to change to **MLPRegressor** used also in lab8. This is a module supplied by scikit learn. It's based on neural network. So with respect to the previous model we lose the “interpretability” of the model but we gain in performance. In fact it's known that NN are kinda of the best in performances. What scikit learn does, is to provide like a wrapper to construct the multi layer perceptron regressor, with just a line of code, and this is an advantage. The disadvantage is that with scikit learn is no possible to use GPUs. In fact GPUs perform way better than CPUs for this kind of matrix computations. Good news is that my account on Jupyter.polito is not enabled to use GPUs, so nothing would be changed. Otherwise, the solution would have been implementing the MLPRegressor from scratch with Pytorch, or there are also some libraries that wraps Pytorch with scikit-learn. Combining the advantages of the two.

For this reason no cross validation was applied for this model. But, different setting of some of the hyperparameters were explored in a different notebook, for avoiding a kernel restart.

I have decided to add **two hidden layers**. For the activation function the **rectified linear unit** has been chosen. This function is defined like:

$$R(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

The biggest advantage of Relu function is that there is no saturation of the gradients. This speed up a lot the stochastic gradient descent, with respect to the tanh and sigmoid functions

**Learning rate:** It is a hyper-parameter that controls how much we are adjusting the weights of our network with respect the loss gradient. The lower the value, the slower we travel along the downward slope. Using a low learning rate might be a good idea, in terms of making sure we do not miss any local minimum, but it can also mean that it will take a lot of time to converge, especially working with CPUs and not GPUs, especially if we get a stuck in sort of plateau region. There is a lot of research in this field how to initialize the weights and how to adjust the learning rate, also because sometimes is difficult to know how big is the “wall” that gets the model stuck in a local minima.

(from B.Caputo *Machine Learning and Deep Learning 2020*).

In this project the **lr** is set to the default value 0.001. Other values has been tried (e.g. 0.01) with the adaptive mode (i.e. the lr is divided by 2 if the loss does not improve for the epochs in `n_iter_no_change`), the classifier gets stuck in a sort of local minimum, so it finds a sub optimal solution, and the performances are worst than setting the initial lr to 0.001 and not adaptive.

**Number of epochs** i.e. `max_iteration` is set to 180, because after that there no significant improvement on the loss.

```
[27]: clf = MLPRegressor(hidden_layer_sizes=(128,64), activation='relu',
    ↪ verbose=True, solver='sgd', max_iter=180, tol=0.01, n_iter_no_change=25)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_val)
y_true = y_val
```



```
r2 = r2_score(y_true, y_pred)
print("r2 score for MLPRegressor: {}".format(r2))
```

```
Iteration 1, loss = 58.47685619
Iteration 2, loss = 25.16348556
Iteration 3, loss = 21.35660277
Iteration 4, loss = 19.19306005
Iteration 5, loss = 17.24127753
Iteration 6, loss = 15.58359859
Iteration 7, loss = 14.91037985
Iteration 8, loss = 13.53853254
Iteration 9, loss = 12.93454563
Iteration 10, loss = 12.24097862
Iteration 11, loss = 11.43025652
Iteration 12, loss = 10.72454331
Iteration 13, loss = 10.23076475
Iteration 14, loss = 9.63218890
Iteration 15, loss = 9.31063889
Iteration 16, loss = 8.60000682
Iteration 17, loss = 8.32127603
Iteration 18, loss = 7.69016729
Iteration 19, loss = 7.31025141
Iteration 20, loss = 6.92619066
Iteration 21, loss = 6.46324693
Iteration 22, loss = 6.36946246
Iteration 23, loss = 5.74046626
Iteration 24, loss = 5.46369452
Iteration 25, loss = 5.32415754
Iteration 26, loss = 4.85906767
Iteration 27, loss = 4.73114508
Iteration 28, loss = 4.20550025
Iteration 29, loss = 4.17765840
Iteration 30, loss = 3.67226107
Iteration 31, loss = 3.56727229
Iteration 32, loss = 3.28176759
Iteration 33, loss = 3.04264505
Iteration 34, loss = 2.89243722
Iteration 35, loss = 2.84081063
Iteration 36, loss = 2.66708737
Iteration 37, loss = 2.44798872
Iteration 38, loss = 2.31497093
Iteration 39, loss = 2.09649200
Iteration 40, loss = 2.08805370
Iteration 41, loss = 1.84330488
Iteration 42, loss = 1.84571495
Iteration 43, loss = 1.71104580
Iteration 44, loss = 1.60835695
```

Iteration 45, loss = 1.54390828  
Iteration 46, loss = 1.37765020  
Iteration 47, loss = 1.30599617  
Iteration 48, loss = 1.23521475  
Iteration 49, loss = 1.20279614  
Iteration 50, loss = 1.10304504  
Iteration 51, loss = 1.02341135  
Iteration 52, loss = 0.96931766  
Iteration 53, loss = 0.94664730  
Iteration 54, loss = 0.80740347  
Iteration 55, loss = 0.90174755  
Iteration 56, loss = 0.83985726  
Iteration 57, loss = 0.71124345  
Iteration 58, loss = 0.66794968  
Iteration 59, loss = 0.64766510  
Iteration 60, loss = 0.62313130  
Iteration 61, loss = 0.53035063  
Iteration 62, loss = 0.61238340  
Iteration 63, loss = 0.50913932  
Iteration 64, loss = 0.55860235  
Iteration 65, loss = 0.53516845  
Iteration 66, loss = 0.47845537  
Iteration 67, loss = 0.39624716  
Iteration 68, loss = 0.39875819  
Iteration 69, loss = 0.40593551  
Iteration 70, loss = 0.39456990  
Iteration 71, loss = 0.34928932  
Iteration 72, loss = 0.34412401  
Iteration 73, loss = 0.32945828  
Iteration 74, loss = 0.38896205  
Iteration 75, loss = 0.26433318  
Iteration 76, loss = 0.30199256  
Iteration 77, loss = 0.25010917  
Iteration 78, loss = 0.26085895  
Iteration 79, loss = 0.26162013  
Iteration 80, loss = 0.26463435  
Iteration 81, loss = 0.21590483  
Iteration 82, loss = 0.28255589  
Iteration 83, loss = 0.21021971  
Iteration 84, loss = 0.21090807  
Iteration 85, loss = 0.19127355  
Iteration 86, loss = 0.24351502  
Iteration 87, loss = 0.22158106  
Iteration 88, loss = 0.13632376  
Iteration 89, loss = 0.13395289  
Iteration 90, loss = 0.14869901  
Iteration 91, loss = 0.14590884  
Iteration 92, loss = 0.14121220

Iteration 93, loss = 0.18210923  
Iteration 94, loss = 0.14424816  
Iteration 95, loss = 0.11852419  
Iteration 96, loss = 0.09728541  
Iteration 97, loss = 0.12577234  
Iteration 98, loss = 0.12344787  
Iteration 99, loss = 0.07134005  
Iteration 100, loss = 0.08569594  
Iteration 101, loss = 0.07893226  
Iteration 102, loss = 0.10980140  
Iteration 103, loss = 0.08785908  
Iteration 104, loss = 0.11183543  
Iteration 105, loss = 0.07651750  
Iteration 106, loss = 0.12581083  
Iteration 107, loss = 0.09440404  
Iteration 108, loss = 0.09167095  
Iteration 109, loss = 0.05456321  
Iteration 110, loss = 0.08009692  
Iteration 111, loss = 0.08122309  
Iteration 112, loss = 0.05742525  
Iteration 113, loss = 0.03801444  
Iteration 114, loss = 0.07045892  
Iteration 115, loss = 0.10070407  
Iteration 116, loss = 0.03543064  
Iteration 117, loss = 0.05464291  
Iteration 118, loss = 0.02947829  
Iteration 119, loss = 0.02396475  
Iteration 120, loss = 0.03584108  
Iteration 121, loss = 0.04067001  
Iteration 122, loss = 0.09489499  
Iteration 123, loss = 0.09842790  
Iteration 124, loss = 0.06235426  
Iteration 125, loss = 0.03358241  
Iteration 126, loss = 0.05644204  
Iteration 127, loss = 0.02853013  
Iteration 128, loss = 0.03588636  
Iteration 129, loss = 0.04153127  
Iteration 130, loss = 0.03996616  
Iteration 131, loss = 0.02651245  
Iteration 132, loss = 0.02618720  
Iteration 133, loss = 0.04347013  
Iteration 134, loss = 0.04665164  
Iteration 135, loss = 0.02895753  
Iteration 136, loss = 0.02792094  
Iteration 137, loss = 0.03778098  
Iteration 138, loss = 0.07660478  
Iteration 139, loss = 0.02431145  
Training loss did not improve more than tol=0.010000 for 25 consecutive epochs.

Stopping.

r2 score for MLPRegressor: 0.7966121099596044

```
[31]: clf.fit(X_dev, y_dev)
      y_pred = clf.predict(X_ev)

      pred_df = pd.DataFrame()
      pred_df['Id'] = wine_df_ev.index
      pred_df['Predicted'] = y_pred

      pred_df.to_csv('Final_MLP_128_64_it180_25.csv', header=True, index=None)
```

```
Iteration 1, loss = 53.36497278
Iteration 2, loss = 23.74714711
Iteration 3, loss = 20.20465267
Iteration 4, loss = 17.84762389
Iteration 5, loss = 16.19518446
Iteration 6, loss = 15.01944114
Iteration 7, loss = 13.80788894
Iteration 8, loss = 12.87584372
Iteration 9, loss = 12.01514529
Iteration 10, loss = 11.28817093
Iteration 11, loss = 10.55714273
Iteration 12, loss = 10.05196476
Iteration 13, loss = 9.38316831
Iteration 14, loss = 8.80913995
Iteration 15, loss = 8.25002876
Iteration 16, loss = 7.75291178
Iteration 17, loss = 7.23323827
Iteration 18, loss = 6.82079867
Iteration 19, loss = 6.31649001
Iteration 20, loss = 6.17899620
Iteration 21, loss = 5.64048992
Iteration 22, loss = 5.33056768
Iteration 23, loss = 5.12040676
Iteration 24, loss = 4.80308721
Iteration 25, loss = 4.60262889
Iteration 26, loss = 4.12656612
Iteration 27, loss = 3.95758679
Iteration 28, loss = 3.66749813
Iteration 29, loss = 3.39255872
Iteration 30, loss = 3.19532563
Iteration 31, loss = 2.97677860
Iteration 32, loss = 2.99294306
Iteration 33, loss = 2.60290440
Iteration 34, loss = 2.49713324
Iteration 35, loss = 2.39456139
Iteration 36, loss = 2.09070856
```

Iteration 37, loss = 2.13953935  
Iteration 38, loss = 2.00850336  
Iteration 39, loss = 1.87431733  
Iteration 40, loss = 1.73788850  
Iteration 41, loss = 1.60767963  
Iteration 42, loss = 1.57401551  
Iteration 43, loss = 1.40233436  
Iteration 44, loss = 1.34060045  
Iteration 45, loss = 1.28603435  
Iteration 46, loss = 1.23723703  
Iteration 47, loss = 1.14478686  
Iteration 48, loss = 1.11825245  
Iteration 49, loss = 1.03094831  
Iteration 50, loss = 1.00513762  
Iteration 51, loss = 0.97046801  
Iteration 52, loss = 0.87110679  
Iteration 53, loss = 0.85098901  
Iteration 54, loss = 0.79878873  
Iteration 55, loss = 0.71055517  
Iteration 56, loss = 0.61121424  
Iteration 57, loss = 0.65860050  
Iteration 58, loss = 0.65890070  
Iteration 59, loss = 0.67768782  
Iteration 60, loss = 0.56628501  
Iteration 61, loss = 0.55631477  
Iteration 62, loss = 0.52923783  
Iteration 63, loss = 0.47441832  
Iteration 64, loss = 0.55229303  
Iteration 65, loss = 0.43834573  
Iteration 66, loss = 0.43147806  
Iteration 67, loss = 0.43089906  
Iteration 68, loss = 0.37876956  
Iteration 69, loss = 0.45088773  
Iteration 70, loss = 0.31849531  
Iteration 71, loss = 0.34956557  
Iteration 72, loss = 0.34015707  
Iteration 73, loss = 0.26957032  
Iteration 74, loss = 0.32558901  
Iteration 75, loss = 0.27094952  
Iteration 76, loss = 0.25944357  
Iteration 77, loss = 0.31320393  
Iteration 78, loss = 0.24635276  
Iteration 79, loss = 0.25327411  
Iteration 80, loss = 0.26853321  
Iteration 81, loss = 0.18427990  
Iteration 82, loss = 0.22811718  
Iteration 83, loss = 0.23644565  
Iteration 84, loss = 0.15789480

Iteration 85, loss = 0.23140164  
Iteration 86, loss = 0.21176052  
Iteration 87, loss = 0.20074138  
Iteration 88, loss = 0.18211390  
Iteration 89, loss = 0.13745740  
Iteration 90, loss = 0.14058538  
Iteration 91, loss = 0.15662598  
Iteration 92, loss = 0.17309826  
Iteration 93, loss = 0.16437803  
Iteration 94, loss = 0.11623422  
Iteration 95, loss = 0.16177184  
Iteration 96, loss = 0.15461455  
Iteration 97, loss = 0.08285078  
Iteration 98, loss = 0.09560941  
Iteration 99, loss = 0.08151446  
Iteration 100, loss = 0.10533427  
Iteration 101, loss = 0.09414676  
Iteration 102, loss = 0.14858717  
Iteration 103, loss = 0.10624442  
Iteration 104, loss = 0.11834992  
Iteration 105, loss = 0.11927949  
Iteration 106, loss = 0.05708478  
Iteration 107, loss = 0.12087947  
Iteration 108, loss = 0.12409756  
Iteration 109, loss = 0.07621408  
Iteration 110, loss = 0.07016926  
Iteration 111, loss = 0.07578335  
Iteration 112, loss = 0.04543062  
Iteration 113, loss = 0.04400092  
Iteration 114, loss = 0.05110917  
Iteration 115, loss = 0.07250740  
Iteration 116, loss = 0.04395513  
Iteration 117, loss = 0.04282265  
Iteration 118, loss = 0.08876136  
Iteration 119, loss = 0.11124681  
Iteration 120, loss = 0.08364885  
Iteration 121, loss = 0.04589048  
Iteration 122, loss = 0.07778984  
Iteration 123, loss = 0.06918949  
Iteration 124, loss = 0.04321710  
Iteration 125, loss = 0.02956152  
Iteration 126, loss = 0.02385337  
Iteration 127, loss = 0.04695742  
Iteration 128, loss = 0.02987689  
Iteration 129, loss = 0.05024526  
Iteration 130, loss = 0.04113667  
Iteration 131, loss = 0.03713177  
Iteration 132, loss = 0.03088846

```

Iteration 133, loss = 0.03887165
Iteration 134, loss = 0.09398879
Iteration 135, loss = 0.02762025
Iteration 136, loss = 0.03349105
Iteration 137, loss = 0.02952055
Iteration 138, loss = 0.07809873
Iteration 139, loss = 0.03466393
Iteration 140, loss = 0.04561743
Iteration 141, loss = 0.02308116
Iteration 142, loss = 0.01705416
Iteration 143, loss = 0.03788704
Iteration 144, loss = 0.02327840
Iteration 145, loss = 0.06512072
Iteration 146, loss = 0.02254824
Iteration 147, loss = 0.01807213
Iteration 148, loss = 0.01649402
Iteration 149, loss = 0.02480563
Iteration 150, loss = 0.02268671
Iteration 151, loss = 0.00972559
Training loss did not improve more than tol=0.010000 for 25 consecutive epochs.
Stopping.

```

```

[33]: clf = MLPRegressor(hidden_layer_sizes=(256,128), activation='relu',
↳ verbose=True, solver='sgd', max_iter=180, tol=0.01, n_iter_no_change=35)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_val)
y_true = y_val
r2 = r2_score(y_true, y_pred)
print("r2 score for MLPRegressor: {}".format(r2))

```

```

Iteration 1, loss = 58.39732895
Iteration 2, loss = 24.95301632
Iteration 3, loss = 21.08542204
Iteration 4, loss = 18.62028709
Iteration 5, loss = 16.87199302
Iteration 6, loss = 15.41080970
Iteration 7, loss = 14.04697074
Iteration 8, loss = 13.02643491
Iteration 9, loss = 12.32034061
Iteration 10, loss = 11.15985326
Iteration 11, loss = 10.49603342
Iteration 12, loss = 9.69375657
Iteration 13, loss = 8.95657711
Iteration 14, loss = 8.44922600
Iteration 15, loss = 7.81752241
Iteration 16, loss = 7.22114990
Iteration 17, loss = 6.60146457
Iteration 18, loss = 6.10174544

```

Iteration 19, loss = 5.91290717  
Iteration 20, loss = 5.45381228  
Iteration 21, loss = 4.66781691  
Iteration 22, loss = 4.66500247  
Iteration 23, loss = 4.09492473  
Iteration 24, loss = 3.62670287  
Iteration 25, loss = 3.77154771  
Iteration 26, loss = 3.02361906  
Iteration 27, loss = 3.09155965  
Iteration 28, loss = 2.71649758  
Iteration 29, loss = 2.61211096  
Iteration 30, loss = 2.23356877  
Iteration 31, loss = 2.09579770  
Iteration 32, loss = 1.89527027  
Iteration 33, loss = 1.85850085  
Iteration 34, loss = 1.66107625  
Iteration 35, loss = 1.52676013  
Iteration 36, loss = 1.34920082  
Iteration 37, loss = 1.23627939  
Iteration 38, loss = 1.10508870  
Iteration 39, loss = 0.92529400  
Iteration 40, loss = 0.95863039  
Iteration 41, loss = 0.92028512  
Iteration 42, loss = 0.78665028  
Iteration 43, loss = 0.77284990  
Iteration 44, loss = 0.67728648  
Iteration 45, loss = 0.61660715  
Iteration 46, loss = 0.60579218  
Iteration 48, loss = 0.47688840  
Iteration 49, loss = 0.50901907  
Iteration 50, loss = 0.40758613  
Iteration 51, loss = 0.34591803  
Iteration 52, loss = 0.44444928  
Iteration 53, loss = 0.29357474  
Iteration 54, loss = 0.34230339  
Iteration 55, loss = 0.32929436  
Iteration 56, loss = 0.28695597  
Iteration 57, loss = 0.22179391  
Iteration 58, loss = 0.27113030  
Iteration 59, loss = 0.27768308  
Iteration 60, loss = 0.20947748  
Iteration 61, loss = 0.12855193  
Iteration 62, loss = 0.16904783  
Iteration 63, loss = 0.15769236  
Iteration 64, loss = 0.11590922  
Iteration 65, loss = 0.12900534  
Iteration 66, loss = 0.13412558  
Iteration 67, loss = 0.12245052



Iteration 68, loss = 0.09282799  
Iteration 69, loss = 0.11977312  
Iteration 70, loss = 0.06666080  
Iteration 71, loss = 0.09025770  
Iteration 72, loss = 0.10038424  
Iteration 73, loss = 0.09566832  
Iteration 74, loss = 0.07666478  
Iteration 75, loss = 0.05978297  
Iteration 76, loss = 0.08507549  
Iteration 77, loss = 0.06519219  
Iteration 78, loss = 0.04885032  
Iteration 79, loss = 0.06880033  
Iteration 80, loss = 0.05298001  
Iteration 81, loss = 0.05457982  
Iteration 82, loss = 0.03630402  
Iteration 83, loss = 0.02582118  
Iteration 84, loss = 0.03333586  
Iteration 85, loss = 0.03759142  
Iteration 86, loss = 0.02400483  
Iteration 87, loss = 0.01999782  
Iteration 88, loss = 0.03532774  
Iteration 89, loss = 0.02538304  
Iteration 90, loss = 0.05575188  
Iteration 91, loss = 0.03149397  
Iteration 92, loss = 0.04383536  
Iteration 93, loss = 0.01573186  
Iteration 94, loss = 0.01233173  
Iteration 95, loss = 0.02157185  
Iteration 96, loss = 0.01309035  
Iteration 97, loss = 0.01869422  
Iteration 98, loss = 0.00912955  
Iteration 99, loss = 0.00660146  
Iteration 100, loss = 0.02660913  
Iteration 101, loss = 0.01106242  
Iteration 102, loss = 0.00671042  
Iteration 103, loss = 0.00872749  
Iteration 104, loss = 0.01531431  
Iteration 105, loss = 0.01143603  
Iteration 106, loss = 0.00578916  
Iteration 107, loss = 0.00388242  
Iteration 108, loss = 0.00331832  
Iteration 109, loss = 0.01524782  
Iteration 110, loss = 0.00724405  
Iteration 111, loss = 0.00689096  
Iteration 112, loss = 0.00346474  
Iteration 113, loss = 0.00406279  
Iteration 114, loss = 0.00857455  
Iteration 115, loss = 0.01309257

```
Iteration 116, loss = 0.02126272
Iteration 117, loss = 0.00831047
Iteration 118, loss = 0.00848713
Iteration 119, loss = 0.00530886
Training loss did not improve more than tol=0.010000 for 35 consecutive epochs.
Stopping.
r2 score for MLPRegressor: 0.8107273215720792
```

```
[35]: clf.fit(X_dev, y_dev)
      y_pred = clf.predict(X_ev)

      pred_df = pd.DataFrame()
      pred_df['Id'] = wine_df_ev.index
      pred_df['Predicted'] = y_pred

      pred_df.to_csv('Final_MLP_256_6128_it180_35.csv', header=True, index=None)
```

```
Iteration 1, loss = 50.80821905
Iteration 2, loss = 23.64980095
Iteration 3, loss = 19.87913044
Iteration 4, loss = 17.64661621
Iteration 5, loss = 15.91735639
Iteration 6, loss = 14.39996418
Iteration 7, loss = 13.23667957
Iteration 8, loss = 12.15835192
Iteration 9, loss = 11.20425640
Iteration 10, loss = 10.24779404
Iteration 11, loss = 9.69625928
Iteration 12, loss = 9.02875421
Iteration 13, loss = 7.98453979
Iteration 14, loss = 7.59707924
Iteration 15, loss = 6.87155638
Iteration 16, loss = 6.30777028
Iteration 17, loss = 5.88038006
Iteration 18, loss = 5.27883281
Iteration 19, loss = 4.86689831
Iteration 20, loss = 4.58922743
Iteration 21, loss = 4.21002664
Iteration 22, loss = 3.82264057
Iteration 23, loss = 3.42169762
Iteration 24, loss = 3.08430172
Iteration 25, loss = 3.04422483
Iteration 26, loss = 2.49710747
Iteration 27, loss = 2.57124288
Iteration 28, loss = 2.31565225
Iteration 29, loss = 1.99655903
Iteration 30, loss = 1.90909883
Iteration 31, loss = 1.80410477
```

Iteration 32, loss = 1.56846713  
Iteration 33, loss = 1.39675892  
Iteration 34, loss = 1.30466949  
Iteration 35, loss = 1.32620937  
Iteration 36, loss = 1.12103703  
Iteration 37, loss = 1.06943232  
Iteration 38, loss = 0.93086192  
Iteration 39, loss = 0.83570415  
Iteration 40, loss = 0.85046081  
Iteration 41, loss = 0.78904827  
Iteration 42, loss = 0.62688102  
Iteration 43, loss = 0.63555133  
Iteration 44, loss = 0.62509318  
Iteration 45, loss = 0.54208481  
Iteration 46, loss = 0.52077143  
Iteration 47, loss = 0.42995887  
Iteration 48, loss = 0.40398845  
Iteration 49, loss = 0.44171905  
Iteration 50, loss = 0.37543426  
Iteration 51, loss = 0.38252115  
Iteration 52, loss = 0.31490101  
Iteration 53, loss = 0.24523984  
Iteration 54, loss = 0.21648492  
Iteration 55, loss = 0.26364625  
Iteration 56, loss = 0.23516137  
Iteration 57, loss = 0.24701190  
Iteration 58, loss = 0.16701267  
Iteration 59, loss = 0.19245004  
Iteration 60, loss = 0.18035415  
Iteration 61, loss = 0.15497067  
Iteration 62, loss = 0.15891468  
Iteration 63, loss = 0.13697878  
Iteration 64, loss = 0.15185686  
Iteration 65, loss = 0.10094345  
Iteration 66, loss = 0.16072498  
Iteration 67, loss = 0.07675900  
Iteration 68, loss = 0.09005880  
Iteration 69, loss = 0.09482960  
Iteration 70, loss = 0.11833911  
Iteration 71, loss = 0.11486698  
Iteration 72, loss = 0.08617822  
Iteration 73, loss = 0.06263908  
Iteration 74, loss = 0.05586519  
Iteration 75, loss = 0.05581700  
Iteration 76, loss = 0.07601347  
Iteration 77, loss = 0.04439674  
Iteration 78, loss = 0.04255724  
Iteration 79, loss = 0.04911298

Iteration 80, loss = 0.02566921  
Iteration 81, loss = 0.05461125  
Iteration 82, loss = 0.05266658  
Iteration 83, loss = 0.06815945  
Iteration 84, loss = 0.02581990  
Iteration 85, loss = 0.01987304  
Iteration 86, loss = 0.01543503  
Iteration 87, loss = 0.05064561  
Iteration 88, loss = 0.02549609  
Iteration 89, loss = 0.03389630  
Iteration 90, loss = 0.03137019  
Iteration 91, loss = 0.07446973  
Iteration 92, loss = 0.02577530  
Iteration 93, loss = 0.03165049  
Iteration 94, loss = 0.02993968  
Iteration 95, loss = 0.02045571  
Iteration 96, loss = 0.01256603  
Iteration 97, loss = 0.01242259  
Iteration 98, loss = 0.00879807  
Iteration 99, loss = 0.00717021  
Iteration 100, loss = 0.00639579  
Iteration 101, loss = 0.00826263  
Iteration 102, loss = 0.01686545  
Iteration 103, loss = 0.02168446  
Iteration 104, loss = 0.01444169  
Iteration 105, loss = 0.01008615  
Iteration 106, loss = 0.00622368  
Iteration 107, loss = 0.00424273  
Iteration 108, loss = 0.00650476  
Iteration 109, loss = 0.00457590  
Iteration 110, loss = 0.00938446  
Iteration 111, loss = 0.01380974  
Iteration 112, loss = 0.00829920  
Iteration 113, loss = 0.00610139  
Iteration 114, loss = 0.01230427  
Iteration 115, loss = 0.00873918  
Iteration 116, loss = 0.00760680  
Training loss did not improve more than tol=0.010000 for 35 consecutive epochs.  
Stopping.